

```
# Importing necessary libraries
import pandas as pd

# Load the dataset from the provided path
file_path = 'loan_data.csv'

# Reading the dataset into a pandas dataframe
loan_data = pd.read_csv(file_path)

# Displaying basic information and first few rows of the dataset to
understand its structure
loan_data.info(), loan_data.head()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 396030 entries, 0 to 396029
Data columns (total 27 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   loan_amnt                             396030 non-null  float64
1   term                                  396030 non-null  object
2   int_rate                              396030 non-null  float64
3   installment                           396030 non-null  float64
4   grade                                 396030 non-null  object
5   sub_grade                             396030 non-null  object
6   emp_title                             373103 non-null  object
7   emp_length                            377729 non-null  object
8   home_ownership                        396030 non-null  object
9   annual_inc                            396030 non-null  float64
10  verification_status                   396030 non-null  object
11  issue_d                               396030 non-null  object
12  loan_status                           396030 non-null  object
13  purpose                               396030 non-null  object
14  title                                 394275 non-null  object
15  dti                                    396030 non-null  float64
16  earliest_cr_line                       396030 non-null  object
17  open_acc                               396030 non-null  float64
18  pub_rec                                396030 non-null  float64
19  revol_bal                              396030 non-null  float64
20  revol_util                             395754 non-null  float64
21  total_acc                              396030 non-null  float64
22  initial_list_status                    396030 non-null  object
23  application_type                       396030 non-null  object
24  mort_acc                               358235 non-null  float64
25  pub_rec_bankruptcies                   395495 non-null  float64
26  address                                396030 non-null  object
```

```
dtypes: float64(12), object(15)
```

```
memory usage: 81.6+ MB
```

```
(None,
 loan_amnt      term  int_rate  installment  grade  sub_grade  \
```

0	10000.0	36 months	11.44	329.48	B	B4
1	8000.0	36 months	11.99	265.68	B	B5
2	15600.0	36 months	10.49	506.97	B	B3
3	7200.0	36 months	6.49	220.65	A	A2
4	24375.0	60 months	17.27	609.33	C	C5

	emp_title	emp_length	home_ownership	annual_inc	...
0	Marketing	10+ years	RENT	117000.0	...
1	Credit analyst	4 years	MORTGAGE	65000.0	...
2	Statistician	< 1 year	RENT	43057.0	...
3	Client Advocate	6 years	RENT	54000.0	...
4	Destiny Management Inc.	9 years	MORTGAGE	55000.0	...

	open_acc	pub_rec	revol_bal	revol_util	total_acc
0	16.0	0.0	36369.0	41.8	25.0
1	17.0	0.0	20131.0	53.3	27.0
2	13.0	0.0	11987.0	92.2	26.0
3	6.0	0.0	5472.0	21.5	13.0
4	13.0	0.0	24584.0	69.8	43.0

	application_type	mort_acc	pub_rec_bankruptcies	
0	INDIVIDUAL	0.0	0.0	
1	INDIVIDUAL	3.0	0.0	
2	INDIVIDUAL	0.0	0.0	
3	INDIVIDUAL	0.0	0.0	
4	INDIVIDUAL	1.0	0.0	

	address
0	0174 Michelle Gateway\nMendozaberg, OK 22690
1	1076 Carney Fort Apt. 347\nLoganmouth, SD 05113
2	87025 Mark Dale Apt. 269\nNew Sabrina, WV 05113
3	823 Reid Ford\nDelacruzside, MA 00813
4	679 Luna Roads\nGreggshire, VA 11650

[5 rows x 27 columns])

The dataset contains 396,030 entries and 27 columns, including both numerical and categorical features. The dependent variable for classification will be the loan_status column.

Handling missing values - filling missing numerical columns with the median and categorical columns with mode

```
loan_data['emp_title'].fillna(loan_data['emp_title'].mode()[0],
inplace=True)
loan_data['emp_length'].fillna(loan_data['emp_length'].mode()[0],
inplace=True)
loan_data['title'].fillna(loan_data['title'].mode()[0], inplace=True)
loan_data['revol_util'].fillna(loan_data['revol_util'].median(),
inplace=True)
loan_data['mort_acc'].fillna(loan_data['mort_acc'].median(),
inplace=True)
loan_data['pub_rec_bankruptcies'].fillna(loan_data['pub_rec_bankruptci
es'].median(), inplace=True)
```

Encoding categorical columns using Label Encoding for simplicity (Loan_status and other categories)

```
# List of columns to encode
#cat_columns = ['term', 'grade', 'sub_grade', 'emp_length',
'home_ownership',
                #'verification_status', 'loan_status', 'purpose',
'title', 'initial_list_status',
                #'application_type']
```

Separate categorical and numerical columns based on dtype

```
cat_columns = loan_data.select_dtypes(include=['object',
'category']).columns.tolist()
num_columns = loan_data.select_dtypes(include=['int64',
'float64']).columns.tolist()
cat_columns # displaying columns with categorical values

['term',
'grade',
'sub_grade',
'emp_title',
'emp_length',
'home_ownership',
'verification_status',
'issue_d',
'loan_status',
'purpose',
'title',
```

```
'earliest_cr_line',
'initial_list_status',
'application_type',
'address']
```

num_columns *# displaying columns with numerical values*

```
['loan_amnt',
'int_rate',
'installment',
'annual_inc',
'dti',
'open_acc',
'pub_rec',
'revol_bal',
'revol_util',
'total_acc',
'mort_acc',
'pub_rec_bankruptcies']
```

Initialize the label encoder

```
from sklearn.preprocessing import LabelEncoder
```

```
le = LabelEncoder()
```

Applying label encoding to each of the categorical columns

```
for col in cat_columns:
    loan_data[col] = le.fit_transform(loan_data[col])
```

Now, the dataset should be preprocessed and ready for feature selection and modeling

```
loan_data.info(), loan_data.head()
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 396030 entries, 0 to 396029
```

```
Data columns (total 27 columns):
```

#	Column	Non-Null Count	Dtype
0	loan_amnt	396030 non-null	float64
1	term	396030 non-null	int64
2	int_rate	396030 non-null	float64
3	installment	396030 non-null	float64
4	grade	396030 non-null	int64
5	sub_grade	396030 non-null	int64
6	emp_title	396030 non-null	int64
7	emp_length	396030 non-null	int64
8	home_ownership	396030 non-null	int64
9	annual_inc	396030 non-null	float64
10	verification_status	396030 non-null	int64

11	issue_d	396030	non-null	int64
12	loan_status	396030	non-null	int64
13	purpose	396030	non-null	int64
14	title	396030	non-null	int64
15	dti	396030	non-null	float64
16	earliest_cr_line	396030	non-null	int64
17	open_acc	396030	non-null	float64
18	pub_rec	396030	non-null	float64
19	revol_bal	396030	non-null	float64
20	revol_util	396030	non-null	float64
21	total_acc	396030	non-null	float64
22	initial_list_status	396030	non-null	int64
23	application_type	396030	non-null	int64
24	mort_acc	396030	non-null	float64
25	pub_rec_bankruptcies	396030	non-null	float64
26	address	396030	non-null	int64

dtypes: float64(12), int64(15)

memory usage: 81.6 MB

(None,	loan_amnt	term	int_rate	installment	grade	sub_grade
emp_title \						
0	10000.0	0	11.44	329.48	1	8
80956						
1	8000.0	0	11.99	265.68	1	9
33317						
2	15600.0	0	10.49	506.97	1	7
127182						
3	7200.0	0	6.49	220.65	0	1
27760						
4	24375.0	1	17.27	609.33	2	14
38300						

emp_length	home_ownership	annual_inc	...	open_acc	pub_rec
revol_bal \					
0	1	5	117000.0	...	16.0
36369.0					0.0
1	4	1	65000.0	...	17.0
20131.0					0.0
2	10	5	43057.0	...	13.0
11987.0					0.0
3	6	5	54000.0	...	6.0
5472.0					0.0
4	9	1	55000.0	...	13.0
24584.0					0.0

revol_util	total_acc	initial_list_status	application_type
mort_acc \			
0	41.8	25.0	1
0.0			1

Initializing and training the Decision Tree Classifier

```
tree_clf = DecisionTreeClassifier(random_state=42)
tree_clf.fit(X_train, y_train)

DecisionTreeClassifier(random_state=42)
```

Making predictions on the test set

```
y_pred = tree_clf.predict(X_test)
```

Evaluating model performance

```
accuracy = accuracy_score(y_test, y_pred)
conf_matrix = confusion_matrix(y_test, y_pred)
```

Getting feature importance from the trained model

```
feature_importances = tree_clf.feature_importances_

accuracy
0.7041890765851072

conf_matrix
array([[ 4592, 10985],
       [12445, 51184]])

feature_importances
array([0.04507793, 0.0049767 , 0.0548494 , 0.07461699, 0.00327091,
        0.08698289, 0.04161787, 0.01281924, 0.08427486, 0.01700044,
        0.02013734, 0.05144465, 0.11388186, 0.05735229, 0.00999808,
        0.10329507, 0.09884682, 0.07401184, 0.01049358, 0.00025082,
        0.02847054, 0.00632988])

X.columns
Index(['loan_amnt', 'term', 'int_rate', 'installment', 'grade',
       'sub_grade',
       'emp_length', 'home_ownership', 'annual_inc',
       'verification_status',
       'purpose', 'title', 'dti', 'open_acc', 'pub_rec', 'revol_bal',
       'revol_util', 'total_acc', 'initial_list_status',
       'application_type',
       'mort_acc', 'pub_rec_bankruptcies'],
      dtype='object')

# Create a DataFrame for feature importances
feature_importance= pd.DataFrame({
    'Feature': X.columns,
```

```

    'Importance': feature_importances
})
feature_importance

```

	Feature	Importance
0	loan_amnt	0.045078
1	term	0.004977
2	int_rate	0.054849
3	installment	0.074617
4	grade	0.003271
5	sub_grade	0.086983
6	emp_length	0.041618
7	home_ownership	0.012819
8	annual_inc	0.084275
9	verification_status	0.017000
10	purpose	0.020137
11	title	0.051445
12	dti	0.113882
13	open_acc	0.057352
14	pub_rec	0.009998
15	revol_bal	0.103295
16	revol_util	0.098847
17	total_acc	0.074012
18	initial_list_status	0.010494
19	application_type	0.000251
20	mort_acc	0.028471
21	pub_rec_bankruptcies	0.006330

```

# Sort the DataFrame by importance in descending order
feature_importance = feature_importance.sort_values(by='Importance',
ascending=False).reset_index(drop=True)
feature_importance

```

	Feature	Importance
0	dti	0.113882
1	revol_bal	0.103295
2	revol_util	0.098847
3	sub_grade	0.086983
4	annual_inc	0.084275
5	installment	0.074617
6	total_acc	0.074012
7	open_acc	0.057352
8	int_rate	0.054849
9	title	0.051445
10	loan_amnt	0.045078
11	emp_length	0.041618
12	mort_acc	0.028471
13	purpose	0.020137
14	verification_status	0.017000
15	home_ownership	0.012819


```

16  initial_list_status      0.010494
17          pub_rec         0.009998
18  pub_rec_bankruptcies    0.006330
19          term            0.004977
20          grade           0.003271
21  application_type         0.000251

```

```


import matplotlib.pyplot as plt
import seaborn as sns

```

```

plt.figure(figsize=(8, 6))
sns.heatmap(conf_matrix, annot=True, fmt="d", cmap="Blues")
plt.xlabel('Predicted Label')
plt.ylabel('True Label')
plt.title('Confusion Matrix')
plt.show()

```



Results from Decision Tree Classification:

Model Accuracy: The Decision Tree model achieved an accuracy of 70.42% on the test data, which indicates that the model is correctly predicting the loan_status around 70% of the time.

Confusion Matrix: The confusion matrix indicates the number of correct and incorrect predictions for each class:

```
True Positives (TP): 51,184
True Negatives (TN): 4,592
False Positives (FP): 10,985
False Negatives (FN): 12,445
```

Feature Importance: The Decision Tree model also provides the importance of each feature in determining the loan status.

These features have the most influence in predicting loan status

Tune this model or try other machine learning algorithms to improve the results?

To tune the Decision Tree model, we'll focus on optimizing key hyperparameters to improve performance. The main hyperparameters to tune for Decision Trees are:

```
max_depth: Maximum depth of the tree (controls overfitting).
min_samples_split: Minimum number of samples required to split an
internal node.
min_samples_leaf: Minimum number of samples required to be at a leaf
node.
criterion: The function to measure the quality of a split (e.g.,
"gini" or "entropy").
```

We will use Grid Search Cross-Validation to explore different values of these hyperparameters and find the optimal combination.

```
# Importing the necessary library for hyperparameter tuning
from sklearn.model_selection import GridSearchCV

# Setting up the hyperparameters for tuning
param_grid = {
    'max_depth': [10, 20, 30, None],
    'min_samples_split': [2, 10, 20],
    'min_samples_leaf': [1, 5, 10],
    'criterion': ['gini', 'entropy']
}

# Setting up the GridSearchCV to find the best parameters
grid_search = GridSearchCV(estimator=tree_clf, param_grid=param_grid,
```

```

cv=3, scoring='accuracy', n_jobs=-1,
verbose=1)

# Fitting the grid search to the data
grid_search.fit(X_train, y_train)

# Get the best parameters from the grid search
best_params = grid_search.best_params_

# Train a new Decision Tree with the best parameters
best_tree_clf = grid_search.best_estimator_

# Evaluating the optimized model
y_pred_tuned = best_tree_clf.predict(X_test)
tuned_accuracy = accuracy_score(y_test, y_pred_tuned)
tuned_conf_matrix = confusion_matrix(y_test, y_pred_tuned)

best_params, tuned_accuracy, tuned_conf_matrix


Fitting 3 folds for each of 72 candidates, totalling 216 fits

({'criterion': 'entropy',
  'max_depth': 10,
  'min_samples_leaf': 5,
  'min_samples_split': 20},
0.8033608565007702,
array([[ 1187, 14390],
       [ 1185, 62444]]))

import matplotlib.pyplot as plt
import seaborn as sns

plt.figure(figsize=(8, 6))
sns.heatmap(tuned_conf_matrix, annot=True, fmt="d", cmap="Blues")
plt.xlabel('Predicted Label')
plt.ylabel('True Label')
plt.title('Confusion Matrix')
plt.show()

```



Components Explained GridSearchCV:

This is the main class from Scikit-learn used to perform grid search with cross-validation.
estimator:

estimator=tree_clf: This specifies the machine learning model you want to optimize. In this case, tree_clf is presumably a Decision Tree Classifier (or any other model). This model will be trained with different combinations of hyperparameters as defined in param_grid. param_grid:

param_grid=param_grid: This parameter defines a dictionary where the keys are the hyperparameters of the model, and the values are lists of settings to be tested.

cv:

cv=3: This specifies the number of cross-validation folds. Here, the dataset will be split into 3 parts (folds), and the model will be trained on 2 parts and validated on 1 part in each iteration. This process will repeat for all combinations of parameters, allowing for a more robust evaluation of model performance. scoring:

`scoring='accuracy'`: This specifies the metric to evaluate the performance of the model during cross-validation. In this case, the model will be evaluated based on its accuracy, which is the proportion of correctly predicted instances to the total instances. `n_jobs`:

`n_jobs=-1`: This parameter controls the number of jobs to run in parallel. Setting `n_jobs=-1` means that all available cores will be used, which can significantly speed up the search process. `verbose`:

`verbose=1`: This controls the level of detail displayed during the fitting process. A value of 1 means that progress messages will be printed, allowing you to monitor how the grid search is progressing. Higher values (e.g., 2) provide more detailed messages. **Summary of GridSearchCV Process Exhaustive Search**: GridSearchCV will perform an exhaustive search over all combinations of the parameters specified in `param_grid`.

Cross-Validation: For each combination of parameters, the model will be evaluated using cross-validation. The average accuracy score across all folds will be calculated.

Best Parameters: After evaluating all combinations, GridSearchCV will identify the combination of hyperparameters that resulted in the highest average accuracy.

Final Model: The model is then trained using the best-found parameters on the entire training dataset.

