

Step 1: Importing Libraries and figuring out the data

```
import pandas as pd
```

```
data = pd.read_csv('data.csv')
```

```
data.head()
```

	Unnamed: 0	PassengerId	Survived	Pclass	\
0	0	1	0.0	3	
1	1	2	1.0	1	
2	2	3	1.0	3	
3	3	4	1.0	1	
4	4	5	0.0	3	

	SibSp	\	Name	Sex	Age
0			Braund, Mr. Owen Harris	male	22.0
1					
1	1		Cumings, Mrs. John Bradley (Florence Briggs Th...	female	38.0
1					
2			Heikkinen, Miss. Laina	female	26.0
0					
3			Futrelle, Mrs. Jacques Heath (Lily May Peel)	female	35.0
1					
4			Allen, Mr. William Henry	male	35.0
0					

	Parch	Ticket	Fare	Cabin	Embarked
0	0	A/5 21171	7.2500	NaN	S
1	0	PC 17599	71.2833	C85	C
2	0	STON/O2. 3101282	7.9250	NaN	S
3	0	113803	53.1000	C123	S
4	0	373450	8.0500	NaN	S

```
data.shape
```

```
(1309, 13)
```

```
data.isnull().sum()
```

Unnamed: 0	0
PassengerId	0
Survived	418
Pclass	0
Name	0
Sex	0

```

Age          263
SibSp        0
Parch        0
Ticket       0
Fare         1
Cabin       1014
Embarked     2
dtype: int64

```

```
data.describe()
```

	Unnamed: 0	PassengerId	Survived	Pclass	Age
\count	1309.000000	1309.000000	891.000000	1309.000000	1046.000000
mean	654.000000	655.000000	0.383838	2.294882	29.881138
std	378.020061	378.020061	0.486592	0.837836	14.413493
min	0.000000	1.000000	0.000000	1.000000	0.170000
25%	327.000000	328.000000	0.000000	2.000000	21.000000
50%	654.000000	655.000000	0.000000	3.000000	28.000000
75%	981.000000	982.000000	1.000000	3.000000	39.000000
max	1308.000000	1309.000000	1.000000	3.000000	80.000000

	SibSp	Parch	Fare
count	1309.000000	1309.000000	1308.000000
mean	0.498854	0.385027	33.295479
std	1.041658	0.865560	51.758668
min	0.000000	0.000000	0.000000
25%	0.000000	0.000000	7.895800
50%	0.000000	0.000000	14.454200
75%	1.000000	0.000000	31.275000
max	8.000000	9.000000	512.329200

Step 2: EDA process


2.1 Survival Count Plot: Understand the distribution of the target variable.

```

import seaborn as sns
import matplotlib.pyplot as plt

```


```
# Countplot of Survived
sns.countplot(x='Survived', data=data)
plt.title('Survival Distribution')
plt.show()
```




2.2 Age Distribution: Explore the distribution of Age and its relationship to survival.

```
sns.histplot(data['Age'], kde=True)
plt.title('Age Distribution')
plt.show()


# Relationship between Age and Survival
sns.boxplot(x='Survived', y='Age', data=data)
plt.title('Survival by Age')
plt.show()
```





2.3 Fare Distribution by Survival: Explore how ticket fare affects survival.


```
sns.boxplot(x='Survived', y='Fare', data=data)
plt.title('Survival by Fare')
plt.show()
```




2.4 Categorical Variables vs Survival: For Pclass, Sex, and Embarked.

```
sns.countplot(x='Pclass', hue='Survived', data=data)
plt.title('Survival by Passenger Class')
plt.show()
```

```
sns.countplot(x='Sex', hue='Survived', data=data)
plt.title('Survival by Gender')
plt.show()
```





Step 3: Data Preprocessing

3.1 Handling Missing Values

```
data = data.drop(columns=['Unnamed: 0'])  
  
# Fill missing Age with median  
data['Age'].fillna(data['Age'].median(), inplace=True)  
  
# Fill missing Embarked with mode  
data['Embarked'].fillna(data['Embarked'].mode()[0], inplace=True)  
  
# Fill missing Fare (if any)  
data['Fare'].fillna(data['Fare'].median(), inplace=True)
```

3.2 Encoding Categorical Variables

```
# One-hot encode Embarked  
data = pd.get_dummies(data, columns=['Embarked'], drop_first=True)
```



```
# Encode Sex column
data['Sex'] = data['Sex'].map({'male': 1, 'female': 0})
```

3.3 Feature Scaling

```
# You may want to standardize features like Fare and Age.

from sklearn.preprocessing import StandardScaler

scaler = StandardScaler()
data[['Age', 'Fare']] = scaler.fit_transform(data[['Age', 'Fare']])
```

Step 4: Feature Selection (Lasso Regularization)

Lasso can be used for feature selection by penalizing less important features.

```
# Drop columns that won't be used in the model
data.drop(columns=['Name', 'Ticket', 'Cabin'], inplace=True)
```

```
plt.figure(figsize=(10, 8)) # Set the figure size
sns.heatmap(data, annot=True, cmap='viridis',
            fmt='g') # 'annot' displays values on the heatmap
plt.title('Heatmap of DataFrame')
plt.xlabel('Columns')
plt.ylabel('Rows')
plt.show()
```

```
''' We when checked above knew there are missing values in the target
y(survived) so we
need to deal with it first '''

# Fill missing Survived values with the most frequent value (mode)

data['Survived'].fillna(data['Survived'].mode()[0], inplace=True)

from sklearn.linear_model import LogisticRegression
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.linear_model import LassoCV
import numpy as np

# Split data into X and y
X = data.drop(columns=['Survived'])
y = data['Survived']

# Split into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y,
                                                    test_size=0.2, random_state=42)

# Apply Lasso to perform feature selection
lasso = LassoCV(cv=5, random_state=0).fit(X_train, y_train)
```

```

# Get selected features
coef = np.where(lasso.coef_ != 0)[0]
selected_features = X.columns[coef]
print("Selected features by Lasso:", selected_features)

Selected features by Lasso: Index(['Unnamed: 0', 'PassengerId',
'Pclass', 'Sex'], dtype='object')

# Split data into X and y
X = data[['PassengerId', 'Pclass', 'Sex']]
y = data['Survived']

# Split into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y,
test_size=0.2, random_state=42)

```

Step 5: Model Building

```

from sklearn.linear_model import LogisticRegression

# Create the logistic regression model
model = LogisticRegression()

# Fit the model to the training data
model.fit(X_train, y_train)

# Predicting the results for test set
y_pred = model.predict(X_test)

```

Step 6: Evaluation Metrics

6.1 Accuracy Scores

```

from sklearn.metrics import accuracy_score

accuracy = accuracy_score(y_test, y_pred)
print(f'Accuracy: {accuracy}')

Accuracy: 0.8282442748091603

```

6.2 Precision, Recall, and F1-Score:

```

from sklearn.metrics import precision_score, recall_score, f1_score

precision = precision_score(y_test, y_pred)

```

```
recall = recall_score(y_test, y_pred)
f1 = f1_score(y_test, y_pred)
```

```
print(f'Precision: {precision}')
print(f'Recall: {recall}')
print(f'F1 Score: {f1}')
```

```
Precision: 0.75
Recall: 0.5753424657534246
F1 Score: 0.6511627906976744
```

6.3 Confusion Matrix:

```
from sklearn.metrics import confusion_matrix
```

```
cm = confusion_matrix(y_test, y_pred)
print("Confusion Matrix: \n", cm)
```

Confusion Matrix:

```
[[175  14]
 [ 31  42]]
```


```
import seaborn as sns
import matplotlib.pyplot as plt
from sklearn.metrics import confusion_matrix
```

```
# Compute confusion matrix
cm = confusion_matrix(y_test, y_pred)
```

```
# Plot confusion matrix as a heatmap
plt.figure(figsize=(6,4))
sns.heatmap(cm, annot=True, fmt='d', cmap='Blues', cbar=False)
```

```
# Add labels and titles
plt.title('Confusion Matrix')
plt.xlabel('Predicted Labels')
plt.ylabel('True Labels')
```

```
# Show the plot
plt.show()
```



Step 7: Model Interpretation

```
coefficients = pd.DataFrame({'Feature': X_train.columns,  
                             'Coefficient': model.coef_[0]})  
print(coefficients)
```

	Feature	Coefficient
0	Unnamed: 0	-1.145734
1	PassengerId	1.142903
2	Pclass	-0.806022
3	Sex	-2.118951

```
import statsmodels.api as sm
```

```
X_train_sm = sm.add_constant(X_train)  
logit_model = sm.Logit(y_train, X_train_sm)  
result = logit_model.fit()  
print(result.summary())
```

```
Warning: Maximum number of iterations has been exceeded.  
Current function value: 0.395503  
Iterations: 35
```

Logit Regression Results

```
=====
```

```

Dep. Variable:    Survived    No. Observations:
1047
Model:            Logit      Df Residuals:
1042
Method:           MLE       Df Model:
4
Date:             Tue, 24 Sep 2024    Pseudo R-squ.:
0.3059
Time:             22:44:13    Log-Likelihood:
-414.09
converged:        False    LL-Null:
-596.60
Covariance Type:  nonrobust    LLR p-value:
1.006e-77

```

```

=====
=====
              coef    std err          z      P>|z|      [0.025
0.975]
-----
const         2.3793         nan         nan         nan         nan
nan
Unnamed: 0    -1.2741         nan         nan         nan         nan
nan
PassengerId    1.2712         nan         nan         nan         nan
nan
Pclass        -0.8391     0.104     -8.075     0.000     -1.043
-0.635
Sex           -2.2140     0.185    -11.956     0.000     -2.577
-1.851
=====
=====

```

```

/Users/prakashpandey/Documents/anaconda3/lib/python3.9/site-packages/
statsmodels/base/model.py:607: ConvergenceWarning: Maximum Likelihood
optimization failed to converge. Check mle_retvals
  warnings.warn("Maximum Likelihood optimization failed to "

```

P>|z|: The p-values for Pclass and Sex are both 0.000, which is highly significant, suggesting that these features are statistically significant predictors of survival.

```

# Create a DataFrame for predictions
predictions_df = pd.DataFrame({
    'PassengerId': data.loc[X_test.index, 'PassengerId'], # Use the
indices to match
    'Predicted_Survived': y_pred
})
print(predictions_df)

```

```
# Merge predictions with the original DataFrame
final_df = data.merge(predictions_df, on='PassengerId', how='inner')

# Save the final DataFrame to a CSV file
final_df.to_csv('titanic_with_predictions.csv', index=False)

# Optional: View the columns in the final DataFrame
print(final_df[['PassengerId', 'Survived',
'Predicted_Survived']].head())
```

	PassengerId	Predicted_Survived
1148	1149	0.0
1049	1050	0.0
982	983	0.0
808	809	0.0
1195	1196	0.0
...
572	573	0.0
140	141	1.0
1182	1183	0.0
312	313	1.0
199	200	1.0

[262 rows x 2 columns]

	PassengerId	Survived	Predicted_Survived
0	24	1.0	1.0
1	30	0.0	0.0
2	32	1.0	1.0
3	33	1.0	1.0
4	44	1.0	1.0

```
final_df.tail()
```

	Unnamed: 0	PassengerId	Survived	Pclass	Sex	Age	SibSp	
Parch \								
257	0	1292	1293	0.0	2	1	0.658652	1
258	1	1293	1294	0.0	1	0	-0.581628	0
259	0	1295	1296	0.0	1	1	1.046240	1
260	1	1298	1299	0.0	1	1	1.588862	1
261	0	1301	1302	0.0	3	0	-0.116523	0

	Fare	Embarked_Q	Embarked_S	Predicted_Survived
257	-0.237445	0	1	0.0
258	0.504989	0	0	0.0
259	-0.107504	0	0	0.0

260	3.445726	0	0	0.0
261	-0.493624	1	0	0.0